

MULTI-HORIZON STOCK PRICE FORECASTING USING DEEP LEARNING & RESIDUAL ENSEMBLES

Project Implementation Presentation, Comparative Model Benchmarks & Defense Dossier

Course: Financial Data Analytics (FDA-601)
Target Equity: Apple Inc. (NASDAQ: AAPL)
Data Horizon: Jan 2015 – Jan 2024 (9 Years)

Component: CIA-3 Project Implementation
Institution: CHRIST (Deemed to be University)
Department: AI, ML & Data Science

1. Executive Summary: What Exactly Are We Doing?

Core Objective: Stock price series are notoriously non-stationary, turbulent, and contaminated with noise. Naive deep learning attempts to forecast absolute raw price levels (e.g. directly predicting \$180 from past prices) inevitably suffer from severe lag, trend extrapolation drift, and catastrophic failure when market regimes shift. In this project, we designed, implemented, and rigorously benchmarked a production-grade machine learning system that predicts Apple Inc. (AAPL) next-day stock prices using 9 full years of daily OHLCV trading data (2015–2024).

What We Did:

- Strict Chronological Split:** We divided 2,264 trading days into 70% Train (2015–Aug 2020), 15% Validation (Aug 2020–Jun 2022), and 15% Test (Jun 2022–Jan 2024, exactly 400 held-out trading days). Absolutely zero data leakage occurred.
- 51 Technical Indicators:** We engineered 51 features spanning moving averages, momentum oscillators (RSI, Stochastic, Williams %R), volatility bands (Bollinger, ATR), volume flow (OBV, Chaikin), and cyclical calendar encodings.
- Residual Price Anchoring:** Instead of predicting raw dollars, our networks learn stationary day-over-day price innovations ($\Delta y_t = y_t - y_{t-1}$) and reconstruct price as: $\hat{y}_t = y_{t-1} + \Delta \hat{y}_t$. This is the central architectural innovation.
- Four Model Families:** We trained Ridge Regression, BiLSTM with Bahdanau Attention, TFT-lite Transformer, and a Stacking Ensemble.
- Comprehensive Validation:** Evaluated across 8 quantitative metrics, extracted temporal attention weights, ranked permutation feature importances, and simulated a live long/short algorithmic trading strategy.

2. Detailed Breakdown of the 4 Evaluated Models

Model Name & Full Form	Architectural Type & Internal Mechanics	Strengths & Limitations in Finance
1. Ridge Regression Baseline (L2-Regularized Linear Model)	- Fits linear weights over flattened 60-day × 51-feature input tensor (3,060 dimensions). - Loss: $\min Xw - y ^2 + \lambda w ^2$ with $\lambda=1.0$ selected via 5-fold CV.	Pros: Fast, interpretable, cannot overfit noise. Cons: Incapable of modeling non-linear volatility regimes or temporal interactions.
2. BiLSTM + Attention (Bidirectional Long Short-Term Memory with Bahdanau Additive Attention)	3 Bidirectional LSTM recurrent layers (hidden dims: 128, 64, 32; dropout 0.20). - Processes sequence forward and backward. - Bahdanau attention pools 60 temporal vectors into a single 256-dim context vector.	Pros: Best regression accuracy (RMSE \$3.71, R^2 0.9623). Attention weights reveal memory. Cons: Slower sequential backpropagation than feedforward nets.
3. TFT-lite Transformer (Temporal Fusion Transformer-lite with Multi-Head Self-Attention)	- Linear feature projection (51 → 64 dims) + sinusoidal positional encoding. - Prepends a learnable [CLS] sequence token. 3 Multi-Head Self-Attention layers (4 attention heads, $d_{\text{model}}=64$).	Pros: Captures global all-to-all sequence interactions in parallel; ties BiLSTM (R^2 0.9622). Cons: Requires careful positional encoding to retain temporal order.
4. Stacking Ensemble (Heterogeneous Multi-Branch Meta-Learner)	- Parallel branches: BiLSTM (128-d) + Transformer (64-d) + GRU (64-d). - Latent representations (256-d) feed into an L2 Ridge meta-learner with cross-validated weights.	Pros: Lowest worst-case error (Max Error \$13.47 vs \$14.40). Robust against extreme regime shocks. Cons: Higher compute complexity to train 3 models concurrently.

3. The Mathematical Gamechanger: Residual Price Anchoring

If you ask a neural network to predict raw closing prices directly (e.g., \$175.50), the loss gradients are dominated by the secular trend. When the stock enters a new price territory never seen in training (e.g. rising from \$150 to \$190), naive models fail completely because they have never seen those neuron activation levels.

Formula: $\hat{y}_t = y_{t-1} + \Delta_t$

- y_{t-1} : The known anchor price (yesterday's actual closing price).
- Δ_t : The stationary day-over-day price innovation predicted by the neural network.
- $\hat{y}_t$: The final reconstructed price prediction.

Why it wins top marks: The target innovation Δy has mean ≈ 0 and constant variance. Training on stationary residuals prevents runaway drift, accelerates Huber Loss convergence from 0.02 to 0.00028, and guarantees that predictions remain firmly grounded in recent market reality.

4. Quantitative Results Benchmark & Model Verdict

Model	RMSE (\$) ↓	MAE (\$) ↓	MAPE (%) ↓	SMAPE (%) ↓	R^2 Score ↑	Dir. Acc (%) ↑	Theil's U ↓	Max Err (\$) ↓
Ridge Baseline	5.6738	4.5088	2.91%	2.87%	0.9121	51.38%*	2.1854	\$21.18
BiLSTM + Attention	3.7145*	2.8618*	1.84%*	1.84%*	0.9623*	42.11%	1.4240*	\$14.40
TFT-lite Transformer	3.7195	2.8659	1.84%	1.84%	0.9622	41.85%	1.4256	\$14.55
Stacking Ensemble	4.5219	3.6828	2.30%	2.30%	0.9441	43.36%	1.7409	\$13.47*

VERDICT: Which Model is Best & Why?


- Best Overall Regression Model: BiLSTM + Attention** — Achieves the lowest RMSE (\$3.7145), lowest MAE (\$2.8618), and highest R^2 score (0.9623). Its bidirectional memory gates effectively smooth day-to-day noise while capturing short-term momentum.
- Best Runner-Up: TFT-lite Transformer** — Virtually tied with BiLSTM (R^2 0.9622 vs 0.9623, RMSE \$3.7195). Self-attention heads detect sudden volatility regime shifts with equal fidelity.
- Best for Risk Management: Stacking Ensemble** — Lowest worst-case dollar loss exposure (Max Error \$13.47 vs \$14.40 for BiLSTM). In institutional fund management where tail risk is critical, the Ensemble is safest.

- **The Critical Financial Insight:** While deep learning dominates regression fit ($R^2 > 0.96$), Ridge baseline scores higher on Directional Accuracy (51.38% vs 42.11%). This proves that price magnitude prediction is mathematically distinct from directional movement under market efficiency!

5. Visual Evidence: All 10 Output Plots & Analytical Explanations

Below are all ten publication-quality diagnostic figures generated by our automated pipeline from the held-out 400-day Apple Inc. test set (June 2022 – January 2024). Use these slides and talking points for your presentation.


Figure 1: Training and Validation Huber Loss Convergence Curves



Key Takeaway: Huber Loss converged smoothly from ~0.02 down to 0.00028 across 25 epochs without any divergence or overfitting.

What to Tell Your Teacher: "Here we observe the training and validation loss curves. Notice that both lines decrease monotonically and track each other tightly. Because we used Huber loss with $\delta=1.0$ and Cosine Annealing, the models were protected from earnings outlier shocks, and early stopping terminated training before any overfitting could occur."


Figure 2: Actual vs. Predicted Stock Price Trajectories (400-Day Test Set)



Key Takeaway: BiLSTM+Attention and TFT-lite Transformer closely hug the true AAPL stock trajectory across rapid bull and bear swings.

What to Tell Your Teacher: "In this plot, the dark blue curve is the actual price of Apple stock over 400 held-out trading days, and the colored lines are our model forecasts. Notice how closely the BiLSTM (red) and Transformer (purple) follow major market inflection points, such as the late-2022 dip to \$125 and the 2023 rally to \$195. The Ridge baseline (orange) shows noticeably wider deviation during peak volatility."


Figure 3: Residual Error Distributions and Normal Q-Q Plots



Key Takeaway: Residuals are centered near zero (\$0.00) with a Gaussian bell-curve and linear Q-Q plot alignment, proving unbiased model calibration.


What to Tell Your Teacher: "To verify statistical validity, we examined model residuals. The left histograms show that error is centered right at zero with no systematic over-prediction or under-prediction. The right Q-Q plots show that error quantiles lie cleanly along the normal theoretical line, with only slight fat-tailed deviations at the extremes corresponding to unexpected quarterly earnings gaps."

Figure 4: Comparative Bar Chart of Eight Evaluation Metrics




Key Takeaway: Visual proof that BiLSTM achieves the lowest error (\$3.71 RMSE) and highest explained variance (R^2 0.9623).

What to Tell Your Teacher: "This multi-panel bar chart compares all four models across our evaluation suite. You can see directly that BiLSTM and Transformer dramatically outperform Ridge regression in error metrics (RMSE drops from \$5.67 to \$3.71, MAPE drops from 2.91% to 1.84%), while Ridge maintains a slight edge in Directional Accuracy due to its linear conservatism."

Figure 5: Historical AAPL Price Chart with MAs, RSI, Volume, and Volatility


Key Takeaway: Comprehensive 4-panel historical view of AAPL (2015–2024) highlighting trend regimes, RSI overbought/oversold levels, and volatility spikes.

What to Tell Your Teacher: "This four-panel diagnostic illustrates our dataset foundation over 9 full years. Panel 1 displays Apple's price with 20, 50, and 200-day moving averages and Bollinger Bands. Panel 2 tracks the 14-day RSI oscillator. Panel 3 shows volume surges, and Panel 4 isolates 10-day rolling volatility, clearly highlighting the extreme volatility clusters during the 2020 COVID crash and the 2022 inflation rate hikes."

Figure 6: Pairwise Pearson Correlation Matrix Across Leading Indicators


Key Takeaway: Identifies multicollinear clusters: Moving averages correlate highly (>0.90), while momentum and volume provide independent predictive signals.

What to Tell Your Teacher: "Here is the correlation matrix across our engineered indicators. The deep red square demonstrates that moving average features are highly collinear because they share price origins. However, volume features like OBV and Chaikin Money Flow show near-zero correlation with price levels, which means they inject completely unique, orthogonal information into our neural network."

Figure 7: BiLSTM Bahdanau Temporal Attention Weight Heatmap Across 60 Days


Key Takeaway: Bahdanau attention allocates ~65% of its total weight to the 5 most recent trading days, with a secondary peak at t-20 (monthly cycle).

What to Tell Your Teacher: "This is our model interpretability heatmap. Each row is a test day, and the columns represent the 60-day lookback horizon. Bright yellow indicates high attention. You can clearly see that the model concentrates maximal attention on the most recent 5 trading days (t-1 to t-5), with a distinct secondary band around t-20, confirming that the network independently discovered the 20-day monthly trading cycle!"

Figure 8: Top-20 Features Ranked by Permutation Feature Importance


Key Takeaway: Raw Close price dominates importance (score 0.133), followed by MA trend crossovers (SMA-20/50 cross, MACD signal) and RSI.

What to Tell Your Teacher: "To measure true feature importance without black-box assumptions, we randomly permuted each feature on the test set and measured RMSE degradation. The recent Close price is most critical (0.133), which validates our Residual Price Anchoring formulation. Trend signals like the SMA20/50 cross and MACD signal follow closely, confirming that moving average crossovers provide genuine predictive power."

Figure 9: 20-Day Rolling Mean Absolute Error (MAE) Timeline

Key Takeaway: Demonstrates error regime stability: Deep models keep rolling MAE bounded below \$4 throughout volatile 2022 rate hike cycles.

What to Tell Your Teacher: "A single aggregate RMSE can hide periods of poor performance. Here we plotted the 20-day rolling MAE across time. Notice how the Ridge baseline (blue) spikes up to \$9 error during the turbulent Q4 2022 rate-hike window, whereas our BiLSTM (orange) and Transformer (green) remain remarkably stable with rolling error consistently bounded below \$4."

Figure 10: Algorithmic Long/Short Trading Strategy Cumulative PnL vs. Buy-and-Hold

Key Takeaway: Direct proof of the accuracy-profitability divergence: A 42% Directional Accuracy model cannot beat passive Buy-and-Hold during a strong bull recovery.

What to Tell Your Teacher: "Finally, we backtested an algorithmic trading strategy driven by our models' directional signals (+1 long if predicting an increase, -1 short if predicting a decrease). During the 2022–2024 recovery, the Buy-and-Hold benchmark grew +52.5%. Because our deep models had ~42% directional accuracy, their active short positions dragged performance. This empirically proves the classic quantitative finance lesson: regression accuracy (low RMSE) does not guarantee directional profitability without explicit trading transaction filters!"

6. Presentation Defense & Viva Q&A; Cheat Sheet

Prepare for these exact questions that teachers and evaluators frequently ask:

Q1: Why did you predict price innovation (Δy) instead of predicting the actual stock price?

Answer: Stock prices are mathematically non-stationary with shifting means. If trained directly on raw price levels, neural networks suffer from trend extrapolation drift and lag. Predicting price innovations ($\Delta y = y_t - y_{t-1}$) makes the target stationary (mean ≈ 0). Reconstructing the price as $\hat{y}_t = y_{t-1} + \Delta \hat{y}$ anchors the forecast to today's closing price, eliminating multi-month drift.

Q2: Which model performed the best, and how did you measure it?

Answer: BiLSTM with Bahdanau Attention achieved the best overall regression accuracy with an RMSE of \$3.7145, MAE of \$2.8618, and R^2 of 0.9623. However, the Stacking Ensemble achieved the lowest maximum error (\$13.47 vs \$14.40), making it superior for worst-case risk management, while the Ridge baseline had the highest Directional Accuracy (51.38%).

Q3: What is the difference between your BiLSTM and the Transformer model here?

Answer: BiLSTM processes the 60-day window sequentially in both directions with recurrent gating cells and pools them using additive attention. The TFT-lite Transformer uses all-to-all multi-head self-attention with a [CLS] token and positional encoding, processing the whole sequence in parallel. Both achieved nearly identical R^2 scores (0.9623 vs 0.9622), proving that recurrent and attention mechanisms are equally viable for financial time series at this scale.

Q4: Why did your models have 96% R^2 but only ~42% Directional Accuracy?

Answer: This is the famous 'magnitude vs direction' disconnect in financial econometrics. Our models predict price magnitude with exceptional precision (within \$2.86 on a \$180 stock, or 1.84% MAPE). However, day-to-day noise around the zero line makes sign prediction much harder. A prediction of +\$0.10 when the stock moves -\$0.10 incurs almost zero regression penalty, but is penalized 100% in directional accuracy.

Q5: How did you ensure there was no data leakage in your project?

Answer: We enforced strict chronological splitting without shuffling: Train (first 70%), Validation (next 15%), and Test (final 15%, exactly 400 days). Crucially, the MinMaxScaler was fitted exclusively on the training partition and then applied to transform validation and test sets. No future information ever leaked backward.